Sample answers

The exercises

1. One way of representing an undirected graph G is using an *adjacency matrix*. This is an $n \times n$ matrix, where n is the number of nodes in G. In this matrix, we label the rows and the columns with the names of G's nodes. For a pair (x, y) of nodes, the matrix has a 1 in the (x, y) entry if G has an edge connecting x and y; otherwise the entry has a 0. (As the graph is undirected, this is somewhat redundant: (x, y) and (y, x) will always be identical.) A *complete* graphs is one in which x and y are connected, for all x, y with $x \neq y$. We don't exclude the possibility of a *loop*, that is, an edge connecting a node to itself, although loops are rare.

   How can you tell from its matrix whether an undirected graph

   (a) is complete?
   (b) has a loop?
   (c) has an isolated node?

   Answer.

   (a) All elements (except those on the main diagonal) are 1.
   (b) Some element on the main diagonal is 1.
   (c) Some row (and hence some column) has all zeros.

2. Design an algorithm to check whether two given words are anagrams, that is, whether one can be obtained from the other by permuting its letters. For example, *garner* and *ranger* are anagrams.

   Answer. Note that we cannot just say "for each letter in the first word, check that it occurs in the second, and vice versa." (Why not?)

   A nice solution sorts the letters in each word and simply compares the two results. This solution scales well, that is, has good performance even as the words involved grow longer.

3. Here we consider a function `first_occ`, such that `first_occ`(d, s) returns the first position (in string s) which holds any symbol from the collection d (and returns -1 if there is no such position). For simplicity let us assume both arguments are given as strings. For example, `first_occ("aeiou", s)` means "find the first vowel in `s`." For $s = $ "zzzzzzzzzzzz" this should return -1, and for $s = $ "Phlogiston", it should return 3 (assuming we count from 0).

   Assuming strings are held in arrays, design an algorithm for this and find its complexity as a function of the lengths of d and s.

   Suppose we know that the set of possible symbols that may be used in d has cardinality 256. Can you find a way of utilising this to speed up your algorithm? Hint: Find a way that requires only one scan through d and only one through s.

Answer. Let m be the length of d and let n be the length of s. The obvious algorithm is to scan through s and, for each symbol x met, go through d to see if that was a symbol we were looking for. In the worst case, when no symbols from d are in s, this means making mn comparisons. (We could also structure the search the other way round: for each element of d, find its first occurrence, and then, based on that, decide which occurrence was the earliest of them all; this also leads to mn comparisons in the worst case.)

Knowing that there are 256 possible symbols that may be used in d allows us to introduce an array `first`, indexed by the symbols. (If we are talking ASCII characters then, conveniently, the indices run from 0 to 255.) Initialise this array so all its elements are, say, -1. Now scan through s. For each symbol `x` in s, if `first[x]` is -1, replace it with `x`'s position. Note that we have not considered d at all so far; we have only recorded, in `first`, where the first occurrence of each symbol is (and left the value as -1 for each symbol that isn't in s).

Now all we need is a single scan through d, to find the value of `first[x]` for each symbol `x` in d. As we find these, we keep the smallest such value. It this value is -1, none of the symbols from d was found in s. Otherwise the value is the first occurrence of a symbol from d.

The complexity in this case is $n + m$. For most reasonable values of m and n, this is smaller than mn, so we have a better algorithm, even taking the overhead of maintaining `first` into account. Essentially we have traded in some space (the array `first`) to gain some time. This kind of *time/space trade-off* is a familiar theme in algorithm design.

4. Gaussian elimination, the classical algorithm for solving systems of n linear equations in n unknowns, requires about $\frac{1}{3}n^3$ multiplications, which is the algorithm's basic operation.

   (a) How much longer should we expect Gaussian elimination to spend on 1000 equations, compared with 500 equations?

   (b) You plan to buy a computer that is 1000 times faster than what you currently have. By what factor will the new computer increase the size of systems solvable in the same amount of time as on the old computer?

   Answer.

   (a) The answer is: 8 times longer, and it does not depend on the particular numbers of equations given in the question. We need $\frac{1}{3}(2n)^3$ operations for $2n$ equations and $\frac{1}{3}n^3$ operations for n. The ratio is $\frac{\frac{1}{3}(2n)^3}{\frac{1}{3}n^3} = \frac{8n^3}{n^3} = 8$.

   (b) If the new machine needs t_{new} units of time for a job, the old one needs $t_{old} = 1000\, t_{new}$ units. Let m be the number of equations handled by the new machine in the time the old machine handles n. The new machine handles m in time $\frac{1}{3}m^3$ and n in time $\frac{1}{3}n^3$. So we have
   $$\frac{1}{3}m^3 = 1000 \cdot \frac{1}{3}n^3$$
   Solving for m we get
   $$m = \sqrt[3]{1000\, n^3} = 10n$$
   So we can now solve systems that are 10 times larger.

5. For each of the following pairs of functions, indicate whether the components have the same rate of growth, and if not, which grows faster.

   (a) $n(n+1)$ and $2000\, n^2$ (b) $100n^2$ and $0.01n^3$

   (c) $\log_2 n$ and $\ln n$ (d) $\log_2^2 n$ and $\log_2 n^2$

   (e) 2^{n-1} and 2^n (f) $(n-1)!$ and $n!$

Here ! is the factorial function, and $\log_2^2 n$ is the way we usually write $(\log_2 n)^2$.

Answer.

(a) Same rate of growth.

(b) $0.001n^3$ grows faster than $100n^2$. Namely, as soon as $n > 1$, we have $n^2 < n^3$. So pick the constant c to be 10,000. When $n > 1$, we have $100n^2 < 10,000 \cdot (0.001n^3)$.

(c) Same rate of growth.

(d) $\log_2 n^2 = 2\log_2 n$, so $\log_2 n^2$ has the same rate of growth as $\log_2 n$. However, $\log_2^2 n$ grows faster; namely $\lim_{n\to\infty} \frac{\log_2 n \cdot \log_2 n}{\log_2 n} = \lim_{n\to\infty} \log_2 n = \infty$.

(e) Same rate of growth, since $2^n = 2 \cdot 2^{n-1}$.

(f) $n!$ grows faster than $(n-1)!$; namely $\lim_{n\to\infty} \frac{n!}{(n-1)!} = \lim_{n\to\infty} n = \infty$.

6. (Levitin 2.3.3.) The sample variance of n measurements $x_1, \ldots, x_n$ can be computed as either

$$\frac{\sum_{i=1}^{n}(x_i - \bar{x})^2}{n - 1} \quad \text{where} \quad \bar{x} = \frac{\sum_{i=1}^{n} x_i}{n}$$

or as

$$\frac{\sum_{i=1}^{n} x_i^2 - (\sum_{i=1}^{n} x_i)^2/n}{n - 1}$$

Find and compare the number of divisions, multiplications, and additions/subtractions (additions and subtractions are usually bunched together) that are required for computing the variance according to each of these formulas. For large n, which is better?

Answer. For large n, the second method is better. A bit of counting leads to this table:

	First method	Second method
Divisions	2	2
Multiplications	n	$n + 1$
Additions/subtractions	$3n - 1$	$2n$

7. Use the informal definitions of O, Θ and Ω to determine whether the following assertions are true or false.

(a) $n(n + 1)/2 \in O(n^3)$

(b) $n(n + 1)/2 \in O(n^2)$

(c) $n(n + 1)/2 \in \Theta(n^3)$

(d) $n(n + 1)/2 \in \Omega(n)$

Response

$n(n + 1)/2$ is quadratic, therefore (a) true, (b) true, (c) false, (d) true

8. Prove the running time $T(n) = n^3 + 20n + 1 \quad \in O(n^3)$

Response

By the Big-O definition, $T(n)$ is $O(n^3)$ if $T(n) \leq c.n^3$ for some $n \geq n_0$

Check the condition:

$n^3 + 20n + 1 \leq c.n^3$

then:

$1 + \frac{20}{n^2} + \frac{1}{n^3} \leq c$

Therefore, the Big-O definition holds for $n \geq n_0 = 1$ and $c \geq 22(= 1 + 20 + 1)$

Larger values of n_0 results in smaller factors c (e.g. for $n_o = 10, c \geq 1.201$), but in any case the above statement is valid.

9. Eight balls of equal size are on the table. Seven of them weigh the same, but one is slightly heavier. You have a balance scale that can compare weights. How can you find the heavier ball using only two weighings?

Answer. Put three balls on each weighing pan. If the six balance, use the second weighing to compare the two remaining balls. Otherwise keep just the three balls that were heavier than the three they were compared to. Put one of the three on one pan, and another on the other pan. If one is heavier, that's the heavy ball; otherwise the heavy ball is the one remaining.

10. There are 18 gloves in a drawer: 4 pairs of red gloves, 3 pairs of yellow, and 2 pairs of green. You select gloves in the dark, which means you cannot assess colour nor left/right-handedness. You can check the gloves only after the selection has been made. What is the smallest number of gloves you must pick to have at least one matching pair in the best case? In the worst case?

Answer. The best case is obviously 2 gloves. To analyse the worst case, note that we might pick as many as 4 red, 3 yellow, and 2 green gloves, and yet not have a pair, because for each colour, we happened to take gloves of the same orientation, or "handed-ness". So to be absolutely certain that we have a pair, we need to pick 4+3+2+1 = 10 gloves.

11. After washing 5 distinct pairs of socks, only 8 socks come back from the clothes line. Hence you are left with 4 complete pairs in the best case, and 3 in the worst case. What is the probability of the best case scenario?

Answer. There are $\binom{10}{2} = 45$ ways to select two socks from 10. In five of these cases the two socks have the same colour. These five cases correspond to the best-case scenario, in that we are left with four complete pairs. The probability of this happening is $\frac{5}{45} = \frac{1}{9}$.